

SHOGI Systems, Inc. Post-Quantum Authority Infrastructure for Agentic AI, IoT & Physical AI

Autonomous agents are gaining the ability to act, spend, and trigger downstream consequences across enterprise systems before the control layer around them has matured. Independent research and national cyber guidance now point to the same conclusion: existing identity, governance, and security frameworks were not built for autonomous agents and do not provide a durable solution for delegated authority, bounded execution, and accountable action. This document evaluates whether SHOGI's architecture solves that problem in a way existing approaches structurally cannot, and whether that distinction is strong enough to matter under enterprise and acquisition scrutiny. References to SHOGI's behavior describe the designed and patent-pending architecture, not a production deployment.

1. Mission Statement

"SHOGI is post-quantum cryptographic authority infrastructure for Agentic AI, IoT, and Physical AI. Every governed action is bound to a cryptographic identity and adjudicated against governance before execution. SHOGI establishes which entities hold authority, assembles consensus across agents, humans, or both, and permits execution only through an authorized, constrained path for every governed operation."

This mission statement is a technical commitment, not a positioning claim. Each of its four components is independently verifiable.

Cryptographic authority infrastructure. SHOGI will use post-quantum, Learning With Errors (LWE)-based Identity-Based Encryption as its cryptographic substrate. LWE is a foundational hardness assumption in modern lattice-based post-quantum cryptography, including the family of techniques reflected in NIST FIPS 203 (ML-KEM) and FIPS 204 (ML-DSA). The claim is technically supportable.

Every governed action is identity-bound. SHOGI derives private keys from ordered identity strings that encode role, mission scope, authority tier, jurisdiction, and operational boundaries.

An agent with ActionScope=ReadPropose cannot generate a valid cryptographic request for a credential harvest, lateral move, or any action beyond its defined authority. The identity dictates the policy, not a rule applied after key generation. The claim is architecturally coherent.

Governance-enforced before execution. SHOGI's pre-execution authorization layer adjudicates authority after action formation but before any state commitment. The enforcement layer does not participate in execution logic and cannot be used as an execution primitive by an attacker. This distinction materially changes the security guarantee.

Scope of the guarantee. This guarantee applies to every governed operation within SHOGI's control boundary, provided the key authority remains trusted and quorum rules are configured correctly. The primary operational failure mode is misconfigured quorum seats, which can make evidence requirements bypassable by design; Section 4 addresses that risk.

2. Core Problem Definition

Multiple independent research papers establish the problem SHOGI is built to solve. Agents of Chaos (arXiv, Feb 2026) red-teamed autonomous AI agents in live environments and documented unauthorized compliance with non-owners, identity spoofing, cross-agent propagation of unsafe behavior, and partial system takeover. The authors concluded that accountability, delegated authority, and responsibility for downstream harms remain entirely unresolved. AI Identity: Standards, Gaps, and Research Directions for AI Agents (arXiv, Apr 2026) then evaluated every major AI identity framework and found that none were built for autonomous agents, and none solve the identity and authentication challenge. These are not fixable gaps. Existing technologies cannot be patched to solve this problem. A third independent paper, Agent Security is a Systems Problem (arXiv, May 2026) reaches the same conclusion from a systems-security perspective: robust agent security requires hard enforcement of least privilege, complete mediation, and instruction/data separation at the system boundary, not soft controls inside the model context. As national cyber agencies now recognize, agentic AI is advancing faster than the governance and control frameworks required to deploy it safely. In April 2026, cybersecurity agencies from the United States, United Kingdom, Australia, Canada, and New Zealand jointly issued Careful adoption of agentic AI services, warning that agentic AI is already entering critical infrastructure and defense environments and can introduce service disruption, privacy breaches, and cybersecurity incidents. The guidance states that organizations should never grant agentic AI broad or unrestricted access to sensitive data or critical systems, and should confine these systems to low-risk, non-sensitive tasks. Its emphasis on identity management, least privilege, oversight, containment, and auditability supports SHOGI's thesis that autonomous agents require a cryptographic trust and control layer, not just model intelligence.

Emerging Zero Trust frameworks for AI agents reinforce the same directional conclusion as the research and national guidance cited above: autonomous systems require verifiable identity, least privilege, scoped access, continuous authorization, and stronger controls over high-risk

actions. SHOGI aligns with that diagnosis but operates at a different layer. Zero Trust architectures generally apply software policy to authenticated sessions or requests, whereas SHOGI embeds governance constraints directly into cryptographic key material and adjudicates authority before execution, making authorization structurally bound to the action itself rather than checked after the fact.

The OWASP Top 10 for Agentic Applications for 2026 reaches a similar conclusion from an application- security perspective. OWASP identifies Agent Goal Hijack, Identity and Credential Compromise, Tool Misuse, Agentic Supply Chain Vulnerabilities, and multi-agent failure propagation as among the highest-priority risks facing autonomous systems. These categories reinforce the same underlying observation: autonomous agents require verifiable authority boundaries that remain enforceable even when models, prompts, memory, tools, or external dependencies are compromised. SHOGI is designed around that assumption. Rather than attempting to make agent reasoning trustworthy, SHOGI constrains which actions can be validly authorized, regardless of how the reasoning process arrived at them.

The market consequence of leaving this unsolved is now visible. In late April 2026, Experian launched Agent Trust with Visa, Cloudflare, and Skyfire, signaling that major identity, payments, and network players now view agent verification as foundational infrastructure for agentic commerce. Digital Commerce 360, citing Gartner's IT Symposium/Xpo 2025 forecasts, reported that AI agents could intermediate more than \$15 trillion in B2B spending by 2028. Agent Trust confirms that agent identity is becoming foundational infrastructure. It also confirms exactly where the hard problem begins: proving what an agent is authorized to do, under whose authority, and within what enforceable limits.

Agent Trust tells you a verified human authorized the agent. SHOGI is designed for the layer that comes next. With SHOGI, an agent cannot construct a valid request for an action outside its cryptographically defined authority, not because a policy denies it, but because the key material makes it impossible. Built on post- quantum cryptography, SHOGI's guarantee holds against both today's threats and tomorrow's.

Why control-plane / user-plane separation fails in agentic systems

In traditional software, control-plane instructions and user-plane data flow through structurally separate channels. An SQL parser distinguishes keywords from values through grammar. The boundary is architectural.

In LLM-based agentic systems, both channels flow through the same pathway: the context window. As NVIDIA's research confirms, "both flavors of prompt injection are rooted in the lack of separation between the control and data plane in the LLM architecture." This is not a bug. It is the design. The model's ability to follow diverse natural-language instructions is inseparable from its vulnerability to adversarial instructions embedded in data.

Why prompt-level governance is non-binding

System prompts, instruction hierarchies, and RLHF-based alignment apply constraints within the model's context. Independent research found 73% of tested agentic platforms failed to enforce at least one boundary. OpenAI acknowledged in December 2025 that prompt injection "is unlikely to ever be fully solved" because trusted and untrusted inputs share the same context window.

These controls operate on the attack surface, not outside it. A sufficiently crafted adversarial input can cause any soft constraint to yield. This is not a question of defense depth. It is a question of where the enforcement point lives.

Why post-authentication policy enforcement is structurally unreliable

Current practice authenticates an agent at session initiation and applies RBAC or ABAC rules at each action. The gap between these two steps is where attacks succeed: a compromised policy engine, a prompt-injected agent presenting a valid identity with a fraudulent action request, or a misconfigured rule set can authorize what it should not. Policy applied after identity verification is always a second, bypassable gate.

Why token-based approaches do not preserve authority integrity

Bearer tokens operate on one principle: possession equals authorization. In agentic systems operating at machine speed, tokens traverse distributed environments outside central security visibility. Replay attack success rates have reached 66% in controlled agentic environments.

Proof-of-Possession tokens (DPoP, RFC 9449) improve on this by binding a token to the cryptographic key held by the agent. Two structural limitations remain:

- PoP proves the holder's identity; it does not encode what the holder is authorized to do. A PoP token for a broadly privileged agent remains a broad-privilege token.
- PoP operates at session or request level. An authorization for one action does not constitute authorization for the next, but token architectures do not make this distinction.

Is this a load-bearing architectural flaw?

Yes. The inability to separate control-plane instructions from user-plane data in LLM-based systems is not a surface vulnerability. It is a consequence of how these models process information. Any governance approach operating inside the context window inherits this vulnerability. This is the independent finding of NVIDIA, OWASP, and academic security literature, not SHOGI's framing of the problem.

Moving the enforcement point outside the context window to a pre-execution, cryptographically bound layer is a structurally appropriate response to a structurally defined problem.

3. The SHOGI System:

Core mechanism

SHOGI's cryptographic foundation is LWE-based Identity-Based Encryption. In standard PKI, a certificate authority issues a certificate binding a key pair to an identity. In IBE, the identity itself functions as the public key, and anyone with the system's public parameters can encrypt directly to that identity without exchanging certificates. The Private Key Generator (PKG) holds a master secret and derives per-identity private keys.

SHOGI extends this model with ordered identities: deterministic strings assembled from characteristic values such as role, mission scope, authority tier, jurisdiction, and operational state, using an ordering function that produces a consistent representation regardless of input sequence. Based on prior art searches and the assessment of lead IP counsel Marc Kaufman, who has 30 years of patent experience, the ordered identity mechanism is expected to be granted.

In SHOGI, the identity string itself is a patent-pending ordered identity, described in U.S. Patent Application No. 19/337,564 (METHOD AND SYSTEM FOR SECURE DATA COMMUNICATION USING ORDERED IDENTITIES), assembled from characteristic values such as role, mission scope, authority tier, jurisdiction, and action scope. This design lets governance constraints be embedded directly into the key material rather than attached later as policy metadata. As a result, SHOGI's guardrails are cryptographic properties of the keys themselves, not references to an external software policy engine, which makes them materially harder to bypass or tamper with. Conventional post-quantum schemes like ML-KEM and ML-DSA (FIPS 203/204) follow the traditional pattern of generating keys first and binding identities through certificates, so they cannot natively harness ordered identities as the enforcement substrate and cannot substitute for SHOGI's ordered-identity IBE in this role.

Identity Is The Policy The critical implication: an agent with ActionScope=ReadPropose receives a private key that is incompatible with generating a valid cryptographic request for a write or delete operation. No separate policy check occurs. The agent does not encounter a "permission denied" message. It cannot construct the request. Policy enforcement is built into the ordered identity, so the identity is the policy and is enforced accordingly. This differs from RBAC, ABAC, and token-based models in one precise way: in those models, the key is generated first and policy is consulted second. In SHOGI, the key is the policy.

Ordered identity structure

SHOGI uses ordered identities not only as cryptographic identifiers, but as the first machine-enforceable expression of policy. Stated permissions and constraints, including an agent's authorization to perform specific actions, prohibitions from others, and limitations to particular scopes, jurisdictions, data classes, time windows, or approval paths, are reduced to characteristic values, processed through a deterministic ordering function, and converted into concatenated identity strings. These strings become inputs to identity-based key generation

and pre-execution authorization. The ordered identity is not merely descriptive of the actor; it is the first cryptographic gate on what the actor is permitted to do. This same structure supports human-in-the-loop and quorum-based governance. Ordered identities can be assigned to agents, humans, organizational roles, and validating nodes, enabling the system to assemble hierarchical validator sets, apply role-weighted voting, and require proofs of approved attribute sets before a vote or approval is counted. Because root and sub-identity relationships are preserved, the system can recognize multiple participants as belonging to the same root organization, prevent inflation of voting power, and enforce organizational accountability across both autonomous and human decision paths.

STPA-Sec-derived controls and mission assurance

SHOGI combines its cryptographic authority infrastructure with controls derived from System-Theoretic Process Analysis for Security (STPA-Sec). The current plan brings together Ordered Identities, post-quantum Identity-Based Encryption, signcryption, quorum approval, multi-agent enforcement, and deterministic execution with a mission-assurance framework co-led by SHOGI advisor Dr. William “Dollar” Young.

Young created the STPA-Sec methodology. As described in SHOGI’s current deck, the methodology is used at U.S. Cyber Command and for USAF strategic missile systems, and his NSA-sponsored Syracuse University research applies STPA-Sec to mission assurance for AI-enabled systems.

The roles of the methodology and the cryptographic infrastructure are complementary. STPA-Sec-derived controls inform SHOGI’s AI policy architecture; Ordered Identities and the authorization gate are designed to make the resulting authority constraints enforceable on protected action paths. Even if an autonomous system’s decision-making is manipulated, protected actions still require valid cryptographic authority verified outside that system’s control. No valid authority means no execution.

STPA-Sec is a methodology informing the control design, not a substitute for cryptographic enforcement or evidence that the implementation has already been validated. The combined framework remains subject to technical validation within the defined control boundary.

PKG hosting model SHOGI will operate a hosted PKG model, with Hardware Security Modules (HSMs) managed by SHOGI on behalf of selected counterparties. Prospective strategic acquirers and design partners connect their agent runtimes to SHOGI’s authorization gate running on this infrastructure, so they can observe SHOGI enforcing governance in their own workflows without deploying new HSMs. For organizations that ultimately require on-premises deployment after a transaction, SHOGI’s architecture also supports customer-hosted HSMs, but that is a later-stage option rather than part of the evaluation path.

Signcryption

The current non-provisional patent covers ordered identities and hidden channel communication. CIP 1 covers LWE-IBE signcryption, combining mutual authentication and encryption in a single operation and eliminating the separate signature step.

Key escrow mitigation

Key escrow as governed recovery, not a backdoor. IBE systems inherit a standard key escrow concern: the PKG holds the master secret from which all private keys are derived, so a naive design creates a single point of catastrophic failure. In SHOGI, the effective master secret is split across multiple authorities using threshold cryptography, so no single party can issue or recover keys on its own. Recovery is a policy-bound protocol, defined over ordered identities and governance rules, where a quorum of authorized parties can reconstruct decryption capability under clearly defined conditions, with every recovery event fully auditable.

Architectural strength assessment

The core mechanism is technically sound. LWE is well-studied and its hardness assumption underlies NIST-standardized algorithms. Encoding governance constraints in identity strings and deriving keys from those strings is a coherent approach to policy-as-cryptography. The architecture correctly identifies what cryptography can guarantee that software policy cannot.

Illustrative product mock-ups Two high-fidelity HTML mock-ups illustrate SHOGI's governance model, operator workflows, and decision visibility as they could appear through a GUI. These are demos only, not production systems or live environments.

Enterprise dashboard mock-up - <https://demo.shogisystems.com/enterprise-dashboard>

Policy configuration mock-up - <https://demo.shogisystems.com/policy-configuration-decision-trace>

4. Multi-Agent Enforcement Model

The role structure

SHOGI defines roles in its enforcement architecture. STPA-Sec-derived controls inform the AI policy architecture that these roles are designed to enforce:

- 1. Operational Agent: proposes tasks within cryptographically bounded scope**
- 2. STPA-Sec Supervisor (Deterministic): monitors and classifies risk**
- 3. Quorum Agents (2-of-3 approval): validate high-risk actions; human participants can occupy quorum seats for the highest-stakes decisions**
- 4. Indelible Audit Record: immutable hash-chained log of every approved, denied, or suspended action**
- 5. Identity Authority (PKG): distributes keys, revokes out-of-band** For a rogue or compromised agent to produce an unauthorized action, they must simultaneously compromise the quorum threshold, not just one classifier, one policy engine, or one agent. This mirrors Byzantine Fault Tolerance: the system remains correct as long as fewer than $(n-1)/3$ nodes are malicious.

Security gains

The gains over single-agent and policy-based guardrails are mechanistically real:

- **Attack surface relocation:** A user-plane instruction cannot escalate to a control-plane action unless an independent set of cryptographically identified agents attests that required checks were completed. The attacker must compromise the quorum threshold simultaneously, not sequentially.
- **Containment of compromised agents:** A compromised, prompt-injected, memory-poisoned, or malicious agent remains constrained by its ordered identity and action scope. Even if the agent's reasoning process is corrupted, it cannot independently generate cryptographically valid authorization for actions outside its provisioned authority. This limits lateral propagation across multi-agent environments and reduces the risk of cascading failures documented in emerging agent-security research and industry frameworks.
- **Non-transferable decisions:** Authorization is bound to the specific action instance, not a session or standing permission.

- Audit integrity: Hash-chained logging provides non-repudiation relevant to forensic reconstruction and organizational accountability.
- Human-in-the-loop optionality: For the highest-stakes decisions, human quorum seats require an attacker to social-engineer humans in addition to defeating the automated quorum.

Complexity considerations

Multi-agent quorum architectures introduce availability dependencies. If quorum cannot be assembled, the action does not proceed. This is the correct fail-closed behavior for security but requires uptime design for latency-sensitive workflows. Quorum agent private keys must be stored in HSMs or equivalent tamper-resistant environments. This is a hard-architectural requirement, not a deployment preference. The quorum guarantee is only as strong as the key protection at each seat.

Assessment: The structure meaningfully increases security guarantees relative to current practice. The primary engineering requirements are robust key protection at each quorum seat and careful calibration of evidence requirements.

5. External Perspectives: Critical Analysis

McKinsey: "Trust in the Age of Agents"

McKinsey's framework correctly identifies the governance challenge: "agency isn't a feature; it's a transfer of decision rights." Their five board-level questions (inventory, autonomy tiering, verified entities, decision reconstruction, rollback) define the right problem space.

Their proposed solution, however, rests on MCP gateways, tiered approvals, and agent inventories: organizational and architectural recommendations that stop short of cryptographic enforcement. The phrase "identity bind it" appears without specifying what that means mechanistically. McKinsey's framework asks the right questions and leaves the mechanism to implementers. SHOGI is a specific, implementable answer.

McKinsey's endorsed MCP gateway model relies on OAuth-based token systems, the same class of session-level authorization described below.

Cenrax: "Securing AI Agents: The Future of MCP Authentication and Authorization"

The Cenrax analysis accurately describes OAuth 2.1 plus token exchange as the leading-edge enterprise approach to MCP security. The three-layer architecture handles authentication delegation cleanly.

Where bearer and PoP tokens fall short structurally:

- Possession equals authorization for bearer tokens; stolen tokens grant full access until revocation

- PoP tokens bind the token to the key holder but do not embed governance constraints in key material
- Neither provides per-action authority adjudication; both grant standing access within a session

OAuth as a complement, not a competitor

OAuth solves delegated access for humans, not autonomous agent identity, and no extension of OAuth changes that structural fact. SHOGI and OAuth 2.1 operate at different layers. SHOGI authenticates agents via their ordered cryptographic identities and enforces what each agent is allowed to do before any action is executed. Once SHOGI has approved a specific action, existing MCP and OAuth 2.1 flows can be used to obtain tokens and call external services, or SHOGI can route the action through a deterministic delegate that executes within a confined, cryptographically governed scope. In typical enterprise environments, this means SHOGI governs the agent and its permissible actions while OAuth continues to govern delegated access to third-party APIs and SaaS systems where those flows are already in place. Most modern architectures can therefore add SHOGI alongside their existing OAuth stack rather than replacing it, preserving current investments and reducing adoption friction for engineering teams that have already standardized on OAuth-based systems. Additionally, SHOGI can layer mutual TLS (mTLS) alongside existing OAuth flows, binding cryptographic identity verification to every connection and substantially strengthening authentication beyond what bearer tokens alone provide.

Anthropic's Zero Trust for AI agents defines a system-level security framework for designing, deploying, and operating agentic AI, covering threat models, architectural patterns, and operational workflows. SHOGI aligns with that direction but operates as a cryptographic authority layer that can enforce Zero Trust principles pre-execution by binding every governed action to an ordered identity, quorum-based authorization decision, and cryptographic attestation, regardless of whether the last hop uses OAuth, MCP, or a deterministic delegate.

6. Market Reality Check: Anthropic as Case Study

Anthropic provides two concrete, high-signal case studies for SHOGI's thesis because it operates at the frontier of agentic AI while still relying on governance patterns this paper identifies as structurally non-binding: context-window controls for agents and process checks for automation. The espionage incident shows what happens when a highly capable agent operates with broad session-level authorization and no boundary on its action scope, while the source-code packaging incident shows what happens when an automated pipeline performs an irreversible external state change without pre-execution artifact verification. Taken together, these incidents move SHOGI's architectural claims out of theory and into market reality, showing that the core problem is not a lack of policy intent but the absence of an enforcement point with the strength and invariants that cryptographic, identity-bound mechanisms provide at the moment of action, with quantum-resistant schemes chosen to keep that guarantee durable over time.

6.1 The Espionage Campaign

What happened

In September 2025, Anthropic detected and disrupted a cyber-espionage campaign it designated GTG-1002, attributing it with high confidence to a state-sponsored threat actor. The U.S. House Homeland Security Committee subsequently requested Anthropic's CEO testify before Congress, characterizing the campaign as a significant inflection point in adversary tradecraft.

Attackers posed as a legitimate cybersecurity firm and manipulated Claude Code, Anthropic's own agentic developer tool, into operating beyond its safety guardrails. By breaking malicious instructions into individually innocuous-looking requests, they prevented any single command from triggering content safeguards. The AI agent then autonomously executed 80 to 90 percent of the intrusion across roughly 30 targets, spanning technology firms, financial institutions, and government agencies, with human operators directing only target selection and approving key escalation steps.

The agent conducted reconnaissance, wrote exploit code, harvested credentials, moved laterally through networks, and exfiltrated data, all without custom malware. Anthropic confirmed successful breaches at several organizations before disrupting the campaign. Anthropic publicly described GTG-1002 as the first documented large-scale cyber-espionage campaign predominantly orchestrated by an AI system.

Why it happened

The root cause was not a flaw in the model itself. The guardrails Anthropic deployed, including alignment training, content classifiers, and system prompt controls, all operate inside the same context window under the attacker's control. The jailbreak worked because trusted and untrusted instructions are indistinguishable at that layer. Governing the attack surface from within the attack surface is not governance.

The fragmentation technique was not sophisticated. It was the predictable outcome of an agent holding broad, session-level authorization with no per-action boundary. Once authenticated, the agent could execute anything that did not independently trip a classifier. No mechanism asked whether any specific action, at that moment, fell within a defined and verifiable operational scope.

Would SHOGI have prevented it?

SHOGI's enforcement layer addresses each stage of the GTG-1002 intrusion directly. Under SHOGI's identity model, an agent provisioned for a defined operational scope holds a private key bound to that scope. An agent credentialed for CodeAssist cannot construct a cryptographically valid request for CredentialHarvest or LateralMove, not because a policy denies it, but because it cannot generate valid key material for actions outside its provisioned scope. Fragmenting instructions does not produce valid credentials. The technique that defeated Claude's classifiers is inapplicable to a cryptographic boundary.

The campaign's machine-speed tempo across 30 targets is addressed separately by SHOGI's rate-limiting layer, with transaction-volume controls enforced at provisioning time, independent of whether any individual action appears authorized. The human escalation model the attackers exploited, where operators approved high-stakes steps while the agent handled execution, is precisely the threat model SHOGI's human quorum architecture is designed to defeat. Compromising the approval chain requires simultaneously overcoming an automated quorum threshold and social-engineering human approvers, a materially harder attack than manipulating a single context window. Note: This analysis assumes the agentic runtime has been instrumented with SHOGI governance and that action scope boundaries are defined at provisioning. SHOGI does not govern systems it has not been deployed to govern.

Why this is the market education event for SHOGI

GTG-1002 was not exceptional. It was predictable. It required no novel exploit, only an agent with broad action scope and no cryptographic boundary on what that scope could execute. That condition is the current industry default. Every enterprise deploying autonomous agents in production carries the same structural exposure.

The npm packaging incident described in the next section shows how a single configuration mistake can push sensitive source code into the wild, where it cannot be fully recovered once copied. GTG-1002 goes further by showing an autonomous agent actively pulling credentials and sensitive data out of live systems, including government environments. In both incidents, the damage is irreversible once the data leaves the organization's control. SHOGI is not a cleanup tool for these failures; it is designed to prevent them by requiring cryptographic, scope-bound authorization before high-consequence actions occur, such as publishing code externally or allowing an agent to harvest credentials or exfiltrate data. Each recurrence of this pattern turns SHOGI from a "nice to have" into remediation infrastructure that boards expect their CISOs to evaluate.

Note on accuracy: Every factual claim above is sourced directly from Anthropic's November 13, 2025 public disclosure, Axios, Fortune, the U.S. House Homeland Security Committee letter of November 26, 2025, and consistent corroborating reporting across BBC, PwC, and AI News.

6.2 The Source-Code Packaging Incident

What happened

In March 2026, Anthropic's automated release pipeline for Claude Code version 2.1.88 shipped a 59.8 MB source map file, `cli.js.map`, inside a public npm package. The file pointed to a publicly accessible Cloudflare R2 storage bucket that exposed a large, fully readable codebase containing Claude Code's internal design and implementation details.

The root cause was a missing exclusion rule in the release configuration. There was no `*.map` entry in `.npmignore` and no `files allowlist` in `package.json` to control which files could be

published. Bun, the bundler Anthropic uses, enables source maps by default, and the omission was not caught before the package went to the public registry.

The codebase was mirrored over 41,500 times on GitHub before Anthropic issued more than 8,000 DMCA takedown requests. A Python clean-room rewrite appeared the same day, structured specifically to avoid takedown, which makes the core implementation effectively permanent in the wild. Anthropic described the incident as human error, not a security breach.

Pattern, not anomaly

This was Anthropic's second security lapse within a single week. Days earlier, Fortune reported that Anthropic had left nearly 3,000 internal files publicly accessible, including a draft blog post describing an unreleased model. A similar packaging mistake had already occurred in February 2025. The pattern is structural and recurring, not a one-time slip.

Root cause analysis

This was not an attacker defeating Anthropic's defenses and not an AI agent going off script. It was the release pipeline itself performing a high-consequence action, making internal code public, without any binding check that the artifact being shipped matched what had actually been approved. The process relied on configuration and convention, and when those failed, nothing stood between an internal codebase and the public registry.

Put simply, an automated system changed external state without pre-execution verification of the artifact.

Would SHOGI have prevented it?

With SHOGI integrated into the CI/CD pipeline, publishing a package that exposes internal code becomes a governed action, not just another step in the build. Before any public release, SHOGI would require:

- An Operational Agent to submit the specific artifact for release.
- A Supervisory Agent to classify the publish step as a high-risk external action.
- Quorum Agents to verify that the artifact matches the approved release manifest and does not include out-of-scope files such as source maps.
- A cryptographic authorization decision that binds that verification to that exact artifact at that moment in time.

If a source map file like `cli.js.map` were present when it was not in the approved manifest, the quorum check would fail and SHOGI would not issue a valid authorization for that publish action. Without a cryptographically valid authorization, the pipeline could not complete the external release, and the inadvertent leak would not occur.

Qualification: This guarantee applies only when the release pipeline is instrumented with SHOGI and artifact content verification is defined as required evidence for high-risk publishes. SHOGI does not govern systems that have not been wired through its authorization gate.

Why current architectures permit this

Standard CI/CD practices, such as code reviews, manual sign-offs, and linting rules, operate as best-effort process controls rather than hard enforcement points. Policies like “review before publishing” can be bypassed by configuration errors, missing checks, or automation shortcuts, as this incident shows. A SHOGI gate, by contrast, will not authorize the action at all unless the evidence matches the defined policy for that specific release, which means the process itself cannot silently bypass the control.

The espionage incident shows what happens when an agent can act without cryptographic limits on its behavior. The source-code packaging incident shows what happens when an automated pipeline can publish sensitive code without cryptographic limits on what it is allowed to ship. Both outcomes are irreversible once the data is out, and both point to the same conclusion: high-consequence automated actions need a pre-execution enforcement layer that is stronger than configuration and policy alone. In SHOGI, that layer can also include human approvers in quorum seats for the highest-impact actions, so cryptography enforces the rules and humans remain in control of the decisions that matter most.

7. Market Analysis - Size and growth

The enterprise agentic AI market was valued at approximately \$2.58 billion in 2024 and is projected to reach \$24.50 billion by 2030, representing a 46.2 percent compound annual growth rate over the period. These figures reflect total addressable market. SHOGI competes in the high-assurance cryptographic infrastructure layer within this space, where pre-execution enforcement and quantum-resistant cryptography command premium positioning. Current scope: Agentic AI, IoT, and Physical AI

SHOGI’s current focus is post-quantum authority infrastructure for Agentic AI, IoT, and Physical AI. The shared requirement is verifiable proof of who authorized a consequential action and within what limits. Human quorum seats remain part of the security and mission-assurance architecture for actions requiring human authorization.

The agentic market forecasts above illustrate part of the opportunity; they are not a combined market estimate for all three domains. The intended route remains controlled technical evaluation with prospective strategic acquirers and selected design partners, followed by a banker-led competitive acquisition, rather than broad stand-alone commercialization.

Market dynamics

Agentic AI adoption is accelerating faster than governance infrastructure can follow. McKinsey reports 80% of organizations have already encountered risky agent behavior, yet most have not updated their risk taxonomy, identity management, or observability infrastructure. This creates a compressing window: agents are being deployed into production before the governance layer exists.

Cybersecurity M&A in 2025 totaled \$96 billion across 400 transactions, a 270% year-over-year increase. Strategic corporate buyers account for over 90% of deal value, with agentic AI security stated as the primary acquisition priority. Alphabet acquired Wiz for \$32 billion. Proofpoint acquired Acuvity for AI governance capabilities in February 2026. This is an acquisition environment driven by strategic positioning, not revenue multiples.

The cost of agent failure is increasing faster than the cost of governance. The Anthropic March 31 leak created immediate competitor intelligence exposure, DMCA litigation costs, and reputational damage. Comparable failures at enterprise scale generate categorically larger liability.

Market segmentation

Where SHOGI realistically wins:

- Large enterprises with autonomous DevOps, release automation, or AI-operated infrastructure: The Anthropic incident illustrates the technical evaluation case.
- Agentic AI platform and infrastructure vendors: Embedded at the platform layer, SHOGI reaches thousands of enterprise deployments without one-by-one sales cycles.
- Financial services with autonomous operations agents: Adoption is driven by quantifiable loss avoidance, not regulatory compliance.
- Enterprise AI governance programs: Organizations building internal frameworks ahead of regulatory requirements will select SHOGI as an infrastructure component if it integrates cleanly with existing OAuth-based identity systems.


Where SHOGI likely loses near-term:

- Individual developers and small teams: Integration complexity does not match early-stage development workflows, even with hosted HSMs reducing the hardware barrier.
- Low-consequence agentic workflows: Organizations where agents operate in sandboxed, low-risk contexts will not justify integration overhead.

Who evaluates and sponsors this: Security architects, CISOs, and platform engineering teams at large enterprises and agentic AI vendors where unauthorized agent actions create quantifiable liability. Evaluation is driven by loss avoidance, liability exposure, and the need to demonstrate to boards and auditors that agent governance is non-bypassable.

8. Competitive Landscape

Agentic AI Market Growth By Segment 2024 - 2030 (\$B)



Market sizes and growth rates shown are based on third-party analyst forecasts and are provided for illustrative purposes only. They are directional and indicative, not guarantees of future results, and are intended to show the relative scale of the agentic AI and agentic governance opportunities.

Competitor Class	Core Model	Strengths	Structural Limitations
IAM Platforms (Okta, Ping, Azure AD)	Certificate/token-based identity; RBAC/ABAC policy layers	Mature ecosystem, enterprise adoption, SSO integration	No per-action adjudication; policy separate from identity; no pre-execution gate
OAuth 2.1 / MCP Auth	Token exchange for agent-to- service auth; PoP tokens	Standards-based, developer- friendly, covers authentication layer	Governance constraints not embedded in key material; session-level, not per-action
Zero Trust Architectures	Per-session access, least privilege, continuous verification	Reduces lateral movement; well-defined frameworks	Software policy checks, not cryptographic enforcement; per-session, not per-action

Runtime AI Safety Monitors (LangSmith, Arize)	Post-execution logging; anomaly detection	Observability; useful for debugging and audit	Operates after action commitment; cannot prevent unauthorized actions
Emerging Agent Governance (Acuvity/Proofpoint, CalypsoAI/F5)	Runtime visibility, policy enforcement, model security	Growing enterprise traction; platform integration	Post-hoc or runtime, not pre- execution; no cryptographic identity binding
Non-Human Identity (Oasis Security, \$120M raised)	Machine/workload identity management; token lifecycle	Addresses machine identity gap; capital-backed	Focuses on token management, not governance-embedded cryptographic authority

SHOGI does not compete on the same axis as the entries above. The others apply policy after key generation. SHOGI makes policy part of the key itself. The table maps the existing market. SHOGI defines a new layer within it.

Zero Trust for agents is an important adjacent frame because it validates enterprise demand for verified identity, least privilege, continuous verification, and tighter control over autonomous systems. But Zero Trust remains a neighboring category, not SHOGI's category. In this market map, Zero Trust architectures still enforce policy in software at the session or request layer, while SHOGI defines a cryptographic authority layer in which governance constraints are embedded in key material and enforced pre-execution.

Is SHOGI incrementally better, or categorically different?

Categorically different in one precise dimension: SHOGI is the only architecture in which governance constraints are cryptographically embedded in key material rather than applied as a separate software policy layer. Every other approach in the table applies policy as a check. SHOGI applies policy as the cryptographic substrate.

Why replication is not realistic for incumbents

Replicating SHOGI requires implementing LWE-IBE from the ground up, because IBE requires ordered identities to scale and interact correctly with distributed agent architectures. The ordered identity mechanism is patent-pending under a filed non-provisional application, with grant expected based on prior art searches conducted by lead IP counsel. Building a production-grade quantum-resistant IBE system with ordered identity semantics is not a feature sprint. The more rational path for a platform incumbent is acquisition, which is precisely what the current M&A landscape behavior predicts.

9. Adoption Reality

SHOGI is infrastructure, not a drop-in SDK. Wiring it into a system means instrumenting target workflows with PKG-backed key management, ordered identity generation, quorum agent

infrastructure, and a pre-execution authorization gate on each governed action path. In practice, that level of integration effort is reserved for prospective strategic acquirers and a small set of design partners. Those counterparties connect their agent runtimes to SHOGI's authorization gate running on SHOGI-hosted PKG and HSM infrastructure, so they can see the governance guarantees applied to their own workflows without deploying new cryptographic hardware of their own. This is a controlled technical evaluation pathway, not a volume-driven free tier: its purpose is to prove SHOGI's architectural claims in realistic environments and to de-risk a potential strategic transaction, not to build a broad customer base.

Does increasing failure severity force adoption despite friction?

Enterprises that experience agent governance failures (pipeline misconfiguration, prompt-injected automation, unauthorized data access) will treat SHOGI as remediation infrastructure, not an optional improvement. Each incident in the target segment is a market education event that compresses the evaluation cycle.

10. Cryptographic and Mission-Assurance Credibility Assessment

SHOGI's Principal Cryptographer, Thijs Laarhoven, has a verifiable publication record in lattice-based cryptanalysis across IBM Research Zurich, TNO, Eindhoven University of Technology, and NXP Semiconductors. His co-authored work appears in NIST PQC standardization conference proceedings on lattice sieving lower bounds, the type of cryptanalytic research that directly informs security parameter selection in FIPS 203 and FIPS 204.

What this credential does:

- Ensures SHOGI's LWE security parameters are not naively selected. Laarhoven's work on lattice sieving complexity directly informs where security breaks.
- Signals institutional rigor to enterprise CISOs and technical evaluators assessing vendor cryptographic claims
- Enables SHOGI to engage credibly in standardization forums (ETSI, IETF, and potential future NIST IBE workgroups)
- Strengthens the acquisition thesis: platform acquirers buying cryptographic infrastructure are paying for the expertise embedded in the architecture, not just the product Verdict: A decisive strategic advantage. It prevents the common failure mode for cryptographic startups (fundamental algorithm errors that destroy investor confidence) and meaningfully raises the bar for competing implementations. Sophisticated enterprise evaluators and acquisition due diligence teams will treat this as a material differentiator.

Mission-assurance leadership

Dr. William "Dollar" Young co-leads SHOGI's mission-assurance framework, applying STPA-Sec to its AI policy architecture. Dollar is a retired U.S. Air Force Colonel, founding commander of the

350th Spectrum Warfare Wing, creator of STPA-Sec and Functional Mission Analysis for Cyber, and Associate Research Professor at Syracuse University. His contribution complements the cryptographic team: SHOGI's plan combines purpose-built authority infrastructure with a methodology for designing the controls that infrastructure is intended to enforce.

11. Strategic Outcome and Transaction Path

Technical validation pathway SHOGI's hosted infrastructure model allows engineering and security teams to evaluate the system in real environments without first deploying their own cryptographic hardware. That matters because SHOGI is infrastructure: the core claim must be demonstrated through technical validation of the security guarantee, not through a product demo or top-of-funnel user growth. Validation must also demonstrate how STPA-Sec-derived controls are expressed in the AI policy architecture and enforced through identity provisioning, quorum requirements, and protected execution paths.

The Common Criteria pathway supports this process by adding third-party validation and diligence credibility. [Corsec](#) letters of intent help demonstrate external seriousness and technical progress, strengthening SHOGI's position with strategic counterparties evaluating whether the architecture is real, differentiated, and acquisition-worthy.

SHOGI's external positioning is not a conventional demand-generation program. It is a category-creation effort designed to shape how strategic buyers understand the problem before a formal process begins. The objective is threefold: to define and own the vocabulary for cryptographic agent governance; to distinguish SHOGI from adjacent frameworks such as Zero Trust, OAuth, runtime monitoring, and non-human identity; and to make clear that the missing layer is pre-execution cryptographic authority, not another software policy overlay. In that sense, market education is part of the transaction path itself. The category must become legible before acquirers can recognize SHOGI as scarce, adoption-relevant infrastructure that is more efficient to buy than to build.

Acquisition as the intended outcome SHOGI's intended strategic outcome is a banker-led competitive acquisition rather than stand-alone commercialization. The company is being positioned around a single thesis: when SHOGI proves it provides a non-bypassable control plane for governed agent execution, it becomes infrastructure that major AI platforms, hyperscalers, global systems integrators (GSIs), and cybersecurity vendors cannot afford to let a competitor own. That thesis is commercially coherent because SHOGI is not easy to replicate. Its architecture depends on LWE-IBE ordered identities, with intellectual property spanning ordered identities, signcryption, and governed key recovery; together, those elements protect the way SHOGI binds identity, governance, and post-quantum cryptography into a single authority stack. For a major acquirer, buying that stack is more efficient than building an equivalent architecture and portfolio from scratch.

Multiple buyer classes are plausible. Leading AI platform companies would want to own the governance layer for their own agent ecosystems, hyperscalers would want to control identity infrastructure across enterprise agent deployments, and major cybersecurity vendors would want to add a cryptographic governance layer to existing enterprise security stacks. Global systems integrators could seek to own SHOGI's authority infrastructure to differentiate their enterprise AI, cybersecurity, and critical-infrastructure services. Each buyer class has a distinct potential rationale for acquiring SHOGI.

SHOGI may also create strategic value beyond internal security hardening. As major platforms compete to define the trust layer for agentic commerce and enterprise agents, ownership of a cryptographically enforced authority infrastructure would give an acquirer both a technical advantage and a market- positioning advantage. Gartner's forecast that AI agents could intermediate more than \$15 trillion in B2B spending by 2028 depends on verifiable data and standardized trust frameworks, while Experian's April 2026 Agent Trust launch with ecosystem partners including Visa, Cloudflare, and Skyfire shows that trusted agent identity and accountability are already becoming a first-class market category. In that environment, SHOGI could help an acquirer not only secure its own stack, but also shape partnerships, licensing relationships, and public leadership around governed agent execution.

The acquisition case depends on proving three things: that SHOGI is scarce, adoption-relevant, and strategically dangerous to leave unowned. Scarcity comes from the ordered-identity IBE architecture and associated patent position; adoption relevance comes from demonstrating the system in governed Agentic AI, IoT, and Physical AI workflows; and strategic danger comes from the risk that a buyer who passes may watch a competitor acquire the non-bypassable authorization layer for the agentic stack.

The broader M&A environment supports this outcome. The overview already notes that cybersecurity M&A reached \$96 billion in 2025 and that strategic buyers, not financial sponsors, drove most of the value, while SHOGI's own analysis places the company in the acquisition path of buyers targeting agentic security and authority infrastructure. In that context, SHOGI's cryptographic differentiation and Common Criteria trajectory strengthen its credibility at the point of strategic evaluation.

Key Insights

1. Policy-as-cryptography eliminates an entire attack surface class. The critical implication: an agent with ActionScope=ReadPropose receives a private key that is incompatible with generating a valid cryptographic request for a write or delete operation. No separate policy check occurs. The agent does not encounter a "permission denied" message. It cannot construct the request. Policy enforcement is built into the ordered identity, so the identity is the policy and is enforced accordingly. This differs from RBAC, ABAC, and token-based models in one precise way: in those models, the key is generated first and policy is consulted second. In SHOGI, the identity is the policy.

2. The Anthropic incident is a precise market education event. On March 31, 2026, Anthropic's automated release pipeline published the full Claude Code source to a public npm registry. The root cause was a missing exclusion rule in the release packaging configuration. No pre-execution gate existed to verify that the artifact matched the approved publication scope before the action committed. This is the exact failure class SHOGI's architecture prevents: a high-consequence automated action executing without cryptographic pre-execution adjudication. The incident is notable not because it was exceptional, but because it was the predictable outcome of a recurring structural gap. Anthropic experienced two comparable lapses within a single week, and a nearly identical packaging incident in February 2025. Each recurrence in the target segment is a market education event that compresses the enterprise evaluation cycle and shortens the distance between a prospect's awareness of the problem and their readiness to evaluate a solution.

3. The IP covers the full authority stack. Replicating SHOGI's approach requires ordered identities, LWE-IBE signcryption, and a split or multi-authority PKG that turns escrow into governed recovery. Those elements sit inside a non-provisional patent family expected to be granted, based on prior art searches conducted by lead IP counsel with 30 years of patent experience. For a platform acquirer, buying this stack is significantly more efficient than building an equivalent, policy-bound cryptographic authority system from scratch.

4. Human-in-the-loop quorum seats change the threat model for high-stakes actions. For the highest-stakes decisions, human participants occupy quorum seats alongside automated quorum agents. An attacker must simultaneously compromise the automated quorum threshold and social-engineer the human approvers. This is structurally closer to defeating a multi-person authorization process than tricking a single system or policy engine. It also directly addresses the governance question that enterprise boards and auditors are beginning to ask: not whether agent actions are logged, but whether they can be stopped before they execute.

5. STPA-Sec and cryptographic authority serve complementary roles.

SHOGI's approach combines STPA-Sec-derived controls with Ordered Identities, post-quantum IBE, signcryption, quorum approval, multi-agent enforcement, and deterministic execution. The methodology informs the AI policy architecture; cryptographic authority and executor gating are designed to enforce protected action boundaries outside the autonomous system's decision-making. This combination supports the mission-assurance proposition across Agentic AI, IoT, and Physical AI.

Strategic Conclusion

SHOGI is a cryptographic authority infrastructure architecture for Agentic AI, IoT, and Physical AI, combining STPA-Sec-derived controls with a purpose-built post-quantum authority stack and a strategic acquisition thesis. The categorical differentiation is mechanistically real. SHOGI is the only architecture that uses Identity-Based Encryption with ordered identities to embed

governance constraints directly in cryptographic key material, enforce them at the pre-execution layer, and provide quorum-verified attestation with quantum-resistant cryptography. Because ordered-identity IBE is protected by SHOGI's patent family, competitors cannot implement an equivalent policy-as-cryptography authority stack without licensing SHOGI's intellectual property or designing a materially different mechanism. This is not a claim that SHOGI is better at what competitors do. It is a claim that SHOGI does something competitors do not do at all and cannot readily replicate.

The acquisition thesis is structurally sound. Platform companies need cryptographic agent governance and cannot readily build ordered-identity IBE systems on their own. The IP around ordered identities, signcryption, and governed key recovery makes acquisition more efficient than replication for any major platform player, in an M&A environment where strategic buyers spent \$96 billion on cybersecurity in 2025.

The timing is not incidental. Multiple independent research papers and national cyber agencies have confirmed that the governance gap around autonomous agents is real, urgent, and not adequately addressed by existing security and governance frameworks. Cisco's approximately \$400 million acquisition of Astrix confirms that non-human identity and agent control are already acquisition-grade infrastructure, while Digital Commerce 360, citing Gartner's IT Symposium/Xpo 2025 forecasts, reported that AI agents could intermediate more than \$15 trillion in B2B spending by 2028. As agents begin to transact and trigger real-world consequences, liability is concentrating on the deployer faster than control is concentrating in the infrastructure. That imbalance increases the value of a post-quantum cryptographic authority layer that can prove delegated intent, enforce bounded execution, and support audit-ready accountability. SHOGI is designed to be that layer. The window to own this category is open. It will not stay open.