The Authority Layer

How SHOGI compares with twenty existing platforms securing AI agents

November 2, 2026

Investors who hear “AI agent security” often conclude that the space is crowded. We tested that. We reviewed the twenty existing platforms investors are most likely to compare SHOGI with, from early-stage startups to Okta, AWS, Cyera and the agent platform NVIDIA and Anthropic announced in September. For each one we reviewed the information available on their website, the documentation and, where it exists, the public code, and asked one question: what actually stops an agent from doing something it was not allowed to do?

The answer across all platforms was consistent. The agent holds a working credential, or can reach one, and a piece of software decides whether each call may go ahead. Where cryptography is used, it proves who is calling. None of the platforms proves, at the moment of action, that the right people approved that specific action. The space is crowded at the access layer. The authority layer is virtually empty.

SHOGI was created to fill this gap. An agent gets an identity with clear limits, but no power of its own to carry out a critical action. Before any critical action runs, a quorum of independent approvers (agents, and people when the stakes demand it) must sign a one-time permit for that exact action. A separate program that uses no AI, the designated driver, checks that permit and only then carries out the action. SHOGI fits on top of these platforms rather than replacing them, which makes each of them a potential partner, licensee or acquirer.

They authenticate. SHOGI authorizes.

The existing platforms use cryptography like an ID badge: the badge says who the agent is and what it may access, and once through the door the agent can go anywhere the badge allows. SHOGI uses cryptography like a signed work order: one specific job, approved by the right people, traced back to the person who created the agent, checked by the machine before it acts, and valid once. The badge still exists. It is just no longer enough to act.

The timing is not accidental. In May 2026, NSA, CISA and their Five Eyes partners recommended human approval before high-impact agent actions. In October, Cisco’s president called AI agent safety “the issue of our time”. Investors have already put more than $300 million into the smaller startups in this review, and Cyera alone raised $1.4 billion this year. It is now reportedly raising again at a $20 billion valuation, up from $12 billion in June. Buyers are moving: Cisco agreed in May to buy Astrix, a non-human identity company, for a reported $400 million, Rubrik bought Strata in June, and Cyera completed its $1 billion purchase of Oasis Security in September.

Method. Sources are company websites, documentation, public code repositories and press releases, checked between October 19th and 31st, 2026. Funding figures are disclosed amounts only. Ratings reflect what is publicly documented, not what a company may have built privately. The SHOGI row describes its design, not a shipped product. Sources are listed at the end of the annex.

1. The access layer is crowded

Platform	Checks each action before it runs	Agent never holds the credentials	Human approval	Approval by several parties	One-time permit for the exact action	Mission rules built into the identity	Signed proof of each approved action	Post-quantum	In production today	Funding disclosed
Keycard	◐	●	◐	○	○	○	○	○	●	$38M
Arcade.dev	●	●	●	○	○	○	○	○	●	$72M
Okta / Auth0	◐	◐	●	○	○	○	○	○	●	Public
Cyera	●	◐	◐	○	○	○	○	○	●	$1.4B in 2026
AWS Bedrock AgentCore	●	◐	○	○	○	○	○	○	●	Public
NVIDIA OpenShell + Claude	●	●	○	○	○	○	○	○	◐	Public
Outerlimit	◐	◐	○	○	◐	◐	○	○	◐	$16M
AxonFlow	●	◐	●	◐	◐	○	◐	○	●	Not disclosed
Rein Security	●	○	○	○	○	○	○	○	●	$35M
Descope	◐	◐	●	○	○	○	○	○	●	$88M
Permit.io	●	●	●	○	○	○	○	○	●	$14M
Tigera Lynx	●	◐	○	○	○	○	○	○	●	Maker of Calico
Kontext Security	●	○	◐	○	○	○	◐	○	●	$4M
Authrizd	◐	◐	◐	○	○	○	○	○	○	Not disclosed
RuntimeAI	◐	◐	◐	○	○	○	◐	◐	○	Not disclosed
Cerbos	◐	○	○	○	○	○	○	○	●	$11M
AuthZed (SpiceDB)	◐	○	○	○	○	○	○	○	●	$16M
Pomerium	◐	●	○	○	○	○	○	○	◐	$18M
Strata Identity (Maverics)	●	●	◐	○	○	○	○	○	◐	Acquired by Rubrik
Oso	◐	○	○	○	○	○	○	○	●	Assets sold to Tiny
SHOGI	●	●	●	●	●	●	●	●	○	Raising seed

● *yes* ◐ *partly, or claimed without public documentation* ○ *not found in public documentation. Profiles of each platform are in the annex.*

The first three columns are where the rest of the market is focused. Checking each action, keeping credentials away from the agent and pausing for one person are now common practice: Arcade and Permit.io do all three in production. Almost all of the money raised by the startups in this table went into this access layer.

The next five columns are where it is almost empty. Only one platform in this review has started on several approvals and a permit for the exact action: AxonFlow, in an enterprise edition where these pieces were designed but not yet connected. None builds the rules of a mission into the agent's identity, and none is post-quantum at its core. A few others come partway: Kontext signs its local logs, while Outerlimit and RuntimeAI make claims they have not documented.

That is the authority layer, and SHOGI is the only one in this review built for it from the start. SHOGI is at seed stage: one patent application filed, a first demonstrator in development, and a team that has built and certified identity-based encryption products before.

Cyera prices the access layer. AxonFlow proves the authority layer.

Cyera is the best-funded company in this review and the most complete platform at the access layer, valued at $12 billion in June and reportedly raising at $20 billion now. Its Agent Guardian maps what each agent can reach and can block a dangerous tool call inline; through Oasis Security, it checks each request against policy and opens a short-lived session with only the permissions needed. But that is still a session, not a permit for one exact action, with no approval by several parties and no signed proof of who approved what. The market already values the most complete access layer at $20 billion. The authority layer above it, the one even Cyera is missing, is where SHOGI sits.

AxonFlow, a small company that governs AI calls in production, is the closest to SHOGI. Its enterprise edition adds signed decision proofs and approval rules that can require several approvers and exclude the requester. In its public release notes, these pieces were not yet connected to the rest of the product. Its signatures are classical, not post-quantum, and we found no method like STPA-Sec behind its rules. AxonFlow confirms that the authority layer is real. SHOGI's difference is the full chain: rules derived from safety analysis and built into the identity, a permit approved by a quorum, a designated driver that holds the credentials, and post-quantum cryptography throughout.

Why the others cannot simply add it

Any of them can add a second approver: it is a feature, and AxonFlow has started. What a feature cannot add is proof. In their products an approval is a record on their own server, so whoever controls that server or its keys can fake it. Proving an approval takes a key for each approver, a check by the system that acts, for that one action, used once, and credentials moved out of the agent. That is a different architecture, and doing it with post-quantum cryptography is harder still.

2. Authentication is not authorization

Every platform in the table checks something before an agent acts. What matters is what that check rests on.

Take a decision made in software. It protects something only if it cannot be skipped. Kontext's code is public, which makes it a useful example: when its policy engine cannot reach a verdict, because the service is down or the agent's identity is unclear, the call goes through. Developers are never blocked by a broken tool, but the protection disappears exactly when something has gone wrong.

Take a credential. In September, Microsoft described an attack on one of its Azure customers. With two compromised application identities, the attacker tried to delete more than a hundred storage accounts in about seven minutes, and most deletions succeeded. The likely source was a secret posted in a public

GitHub issue. Only the locks protecting the backups held, because the stolen identities had no right to remove them. A compromised identity should never be enough to authorize a destructive action.

What their cryptography proves

The existing platforms use cryptography to authenticate. A signed token says who the agent is, whom it acts for, and what it may access, until the token expires. A token can even describe a narrow grant, such as a single payment. But none of the token-based platforms in this review uses one to require several approvals, carry the rules of a mission, or serve as a one-time permit checked by the machine that acts.

The table below looks at the platforms that use cryptography: what each one signs, what its identity is tied to, and what that proves.

Platform	What it signs or encrypts	What the identity is tied to	What that proves	What it does not prove
Keycard	Access tokens (RS256) carrying the delegation chain.	The agent and the user it acts for.	The token was issued to this agent, for these scopes.	That a specific action was approved. The token covers every action in its scope until it expires.
Okta / Auth0, Descope	OAuth tokens, proof-of-possession (DPoP), cross-app identity assertions.	The agent and the user it acts for.	Who the agent is, and whom it acts for.	That a specific action was approved by the right people.
Permit.io	An Ed25519 signature on the request its gateway forwards.	The gateway.	The request came through the gateway.	Which action was approved, or by whom.
Kontext Security	Optional Ed25519 signatures on a local log of decisions.	The developer's machine.	The local log was not changed after signing.	That nothing was deleted. The log and its key sit on the same machine.
Pomerium	A signed identity assertion (ES256 or RS256).	The user calling.	Who is calling.	What is being done.
Cerbos	Encrypted, signed policy bundles.	No agent identity.	The policy was not altered in transit.	That a decision was made, or obeyed.
Outerlimit	Says credentials are split and rebuilt only when a call is allowed.	Not documented.	Not documented.	No algorithm or proof format has been published.
AxonFlow	Ed25519 signatures on policy bundles and, optionally, on a chain of decision records. Enterprise design: a signed proof for each allowed decision.	The service that signs, with a verified chain from user to agent.	The policy was not altered, and the decision records were not changed.	In public code, that an action was approved before it ran. The proof is enterprise-only and was not yet connected. Not post-quantum.
RuntimeAI	Says it uses ML-KEM and ML-DSA for its key vault and audit log.	An agent key (Ed25519, not post-quantum).	If built as described, that stored keys and logs are intact.	Authority for any action.
SHOGI	A one-time permit for one exact action, signed and encrypted in a single step to the designated driver that will execute it. A signed receipt afterwards.	The person who created the agent, its group and its sub-agents, with the mission's rules built into the identity.	Who authorized which exact action, under which rules, and with whose approval.	That the rules themselves were right. SHOGI proves authority, not good judgment.

Read across, every signature proves who is speaking, or that a record was not changed. None in this review yet proves that a specific action was approved by the right parties, under the rules of a mission, and checked by the machine that carried it out. That is the line between the access layer and the authority layer.

3. The authority layer

The identity carries the rules

Each agent receives an ordered identity, and its key is derived directly from that identity, with no certificate in between. The identity spells out what the agent may do, what it must never do, and what it may do only with approval. In short, the identity is the policy. It is tied to the person who created the agent and to its group, and an agent can hand a sub-agent a narrower identity, never a wider one. Its time window limits how long the key is useful. This is the subject of SHOGI's patent application on Ordered Identities, filed in September 2025 with priority to September 2024, earlier than most of the agent products reviewed here.

The rules come from STPA-Sec, a method that starts from the losses a mission cannot accept and works back to the actions that could cause them. Dr. William "Dollar" Young, its creator, leads SHOGI's STPA-Sec integration. A retired USAF Colonel, he established the 350th Spectrum Warfare Wing and created STPA-Sec, which is used at U.S. Cyber Command and in security engineering for the U.S. Air Force's Sentinel intercontinental ballistic missile program. He participates in the Defense Advanced Research Projects Agency's Information Science and Technology Study Group and is currently developing the protection plan for the Pentagon's AI-assisted classification-management pilot.

Critical actions need a one-time permit

The identity bounds what an agent can ask for. For the actions STPA-Sec marks as critical, that is not enough. A supervisor classifies the risk, a quorum of independent agents must approve the action, and people join the vote when the stakes demand it. SHOGI then issues a permit for that exact action: this target, this time window, used once. Everything else runs without waiting for anyone.

This is the two-person rule the military applies to its most dangerous actions, enforced by cryptography: no agent, and no single person, can authorize a critical action alone. It also goes beyond what NSA, CISA and their Five Eyes partners recommended in May 2026. They asked for human approval before high-impact agent actions; SHOGI makes that approval provable.

A designated driver acts

AI agents can hallucinate or be manipulated, a bit like friends who have had a few drinks. SHOGI gives every protected system a designated driver: a deterministic delegate that runs no AI, has its own identity and holds the real credentials. The permit is signed and encrypted for the driver in a single step, so only the driver can open it and be sure who issued it. The driver then checks the approvals, the time window and that the permit has never been used, performs the exact operation named, and signs a receipt. An auditor given its own identity in the group can verify every receipt.

No one holds the master key alone

In a standard identity-based system, whoever holds the master secret can recreate any key. SHOGI prevents this: the master secret is split among several authorities, and recovering a key requires a quorum of authorized parties, with every recovery recorded. Encryption and signatures are lattice-based, the family behind NIST's post-quantum standards. SHOGI's cryptography is led by Dr. Thijs Laarhoven, its principal cryptographer, whose research underpins those standards. Two further applications, on signcryption and quorum-based key recovery, will be filed after funding.

Speed, repetition and availability

SHOGI only stands in the way of critical actions. Routine and repetitive work runs within the identity's limits without waiting for anyone. Checking a permit is fast: the designated driver verifies it on the spot, with no certificate authority to call. SHOGI is meant for actions that are critical and hard to undo, not for split-second trading. When timing matters, approval can be given ahead of time: the permit names the action and the short window in which it may run. Because approval needs a threshold rather than everyone, one approver being away does not stop the system. If approval cannot be obtained at all, critical actions do not run. SHOGI fails closed, where a platform like Kontext lets calls through.

What it changes

Back to the Azure attack. With SHOGI in place, the stolen secrets would still have let the attacker look around. But every deletion would have needed its own permit, removing backup protection would have needed a human vote, and the designated driver would have refused the rest.

This is the demonstration we are building: an agent whose model has been fully compromised, holding every credential the agent has, still unable to push a single unapproved software release.

SHOGI does not replace these platforms: Okta says who an agent is, Cerbos evaluates rules, OpenShell contains the agent, Arcade routes its calls, and SHOGI gives them the layer they are missing.

Why post-quantum matters now

Today's internet security, from secure websites to digital signatures and crypto wallets, relies heavily on elliptic-curve cryptography. It became the default because it gives the same protection as older methods like RSA with much shorter keys, which makes it fast and well suited to phones, hardware devices, and wallets. It is also commonly used in the first step of a secure web connection, when your browser and a website agree on a shared secret.

That foundation has a known long-term weakness: quantum computers. A large enough quantum computer could break elliptic-curve cryptography outright. None exists today, but security systems take years to replace, so the sensible time to move is before the threat arrives.

Brand-new research adds to the unease. On October 28, 2026, researchers at Ruhr University Bochum and the Max Planck Institute published a paper showing that if a recent AI-produced result from OpenAI holds up, one of the long-held assumptions behind elliptic-curve cryptography is weaker than believed: that no mathematical shortcut exists for breaking it. The OpenAI result, released in mid-October, is still under expert review, and even if it holds, it doesn't make today's encryption any easier to break. But it is another sign that this foundation is less settled than it once seemed.

SHOGI was designed with this in mind from the start. Its identities, action permits, approvals, and execution receipts are built on lattice-based post-quantum cryptography, not the elliptic-curve cryptography used by many existing solutions, because we expected quantum computers to be the long-term threat. This new research wasn't part of that decision, but it reinforces it: SHOGI doesn't rely on the math it puts in question. No cryptography is immune to future breakthroughs, but SHOGI is built on a foundation designed for what comes next.

What an acquirer gains beyond AI agents

The same cryptosystem authorizes any critical command, whether or not an AI issues it: industrial controllers and robots, firmware updates, connected devices and defense systems. Identity-based encryption is already standardized for mission-critical and government communications, through the UK NCSC's MIKEY-SAKKE, ETSI and 3GPP, and ISO/IEC 18033-5. Those schemes are pairing-based and rest on the elliptic-curve mathematics questioned above. ETSI has already evaluated a lattice-based version. SHOGI gives an acquirer a post-quantum successor.

Annex: the twenty existing platforms

Keycard

San Francisco. $38M from a16z, boldstart and Acrew Capital (October 2025). A token service for agents: a Cedar policy runs when a token is issued or exchanged, and an optional gateway swaps in short-lived provider credentials so the agent never sees the real key.

Arcade.dev

San Francisco. $72M, including a $60M Series A led by SYN Ventures with Morgan Stanley and Wipro (June 2026); Cisco Investments joined in September. Every tool call passes through Arcade's engine, which holds the user's OAuth tokens and runs customer-defined checks before calling the API. Sensitive calls can require a push approval. A likely partner rather than a rival: Arcade owns the point where agents act, and SHOGI could be the permit it requires.

Okta / Auth0

The incumbent. Auth0 for AI Agents has been generally available since November 2025, with a token vault and phone-based approval of individual actions. Okta for AI Agents followed in April 2026. Approval is one person; nothing proves afterwards that an action was authorized. Both a rail SHOGI should run on and a logical acquirer.

Cyera

New York and Tel Aviv. Valued at $12 billion in June and reportedly raising at $20 billion in October; raised $1.4 billion in 2026, including $400 million from Goldman Sachs in September. Started in data security. Agent Guardian (August 2026) inventories agents and can block dangerous tool calls inline; Cyera Endpoint extends this to employee devices. Through Oasis Security, bought for $1 billion, it grants short-lived, least-privilege sessions. No multi-party approval, one-time permit or signed proof documented.

AWS Bedrock AgentCore Policy

Generally available since March 2026. A Cedar policy engine attached to AWS's agent gateway checks every tool call against the caller's identity and the call's arguments, denying by default. Dogwood, released in August, adds rules such as "only if an approval happened first". Anything that does not pass through the gateway is not covered, there is no human pause, and decisions are logged in CloudWatch rather than proven. The baseline most investors will compare SHOGI with.

NVIDIA OpenShell + Claude Managed Agents

Announced by NVIDIA and Anthropic in September 2026 and presented by Jensen Huang at GTC Berlin in October. Anthropic runs the agent and keeps its credentials in a vault the agent never sees. NVIDIA's open-source OpenShell runtime checks every action against a deny-by-default policy, logs what it allows or blocks, and uses formal verification to show what an agent can reach. The two are complementary, and OpenShell is a possible integration path.

Outerlimit

London and New York. $16M pre-seed from AlbionVC, Evolution Equity and Crane (September 2026), founded by the team that sold Egress to KnowBe4. Ships an MCP gateway that blocks tools by name and keeps an append-only log. Describes credentials that are split and rebuilt only when a call is allowed, but has published no algorithm and lists credential management as future work. The closest to SHOGI in language.

AxonFlow

Governs AI calls in production; at BukuWarung it governed 1.3 million calls. Public code includes a policy engine, Ed25519-signed policy bundles, optionally signed decision records and holds for human approval. The enterprise edition adds signed decision proofs and approval by several approvers, not yet connected in its public release notes. Its Claude Code plugin lets calls through when the policy service cannot be reached. Classical cryptography. The closest to SHOGI in design.

Rein Security

Tel Aviv and New York. $35M raised, including a $25M Series A (October 2026). A sidecar watches what the agent does at the operating-system level, learns its normal behavior and blocks deviations before they run. The agent keeps its credentials. Rein detects what looks wrong; SHOGI requires proof of what is allowed.

Descope

$88M raised. An identity provider for agents: per-agent OAuth clients, per-tool scopes, a credential vault, and phone-based human approval since June 2026. Policies apply when tokens are issued.

Permit.io

Tel Aviv. $14M raised. Known for OPAL, an open-source policy distribution project. Since March 2026 its MCP Gateway checks every tool call, keeps upstream tokens out of the agent's reach and can pause a call for one administrator to approve. No multi-party approval and no record a third party can verify.

Tigera Lynx

San Francisco. Lynx, generally available since June 2026, governs fleets of AI agents on Kubernetes. It discovers agents, gives each one a cryptographic workload identity, checks traffic between agents, tools and models against Cedar policies, and keeps a fleet-wide audit trail. It governs the fleet; it does not issue a permit for a single action or leave a proof a third party can check.

Kontext Security

Munich. $4M seed (September 2026). Open-source hooks that check each tool call of Claude Code or Codex against a local Cedar policy and can sign a local log of decisions. Covers coding agents on developer machines. Lets calls through when it cannot decide.

Authrizd

San Mateo. One founder, website live since September 2026, no disclosed funding or customers. Its only public artifact is a guided demo labeled synthetic; the grant shown is not signed and a language model interprets the task.

RuntimeAI

San Francisco. One founder, pre-seed. The website claims a very wide range of features, including post-quantum signing, alongside unrelated products. No certification, no public code, documentation behind a verification wall. Its post-quantum vocabulary overlaps SHOGI's; nothing could be tested.

Cerbos

London. $11M raised. A fast, stateless policy engine in Go. Decides; does not enforce, hold credentials or prove. A natural component inside a SHOGI deployment.

AuthZed (SpiceDB)

New York. $16M raised. The leading open-source implementation of Google's Zanzibar model, used by OpenAI. Excellent at answering who may act on what. Blocks nothing by itself.

Pomerium

Los Angeles. $18M raised. An open-source identity-aware proxy. Can block an agent's tool call by tool name and keep upstream credentials away from the agent. Agent support is still marked experimental; no human approval.

Strata Identity (Maverics)

Boulder. Acquired by Rubrik, announced June 2026. Had built a gateway that checks each tool call and exchanges tokens for short-lived, per-tool ones.

Oso

New York. A policy engine. On 6 October 2026 its cloud assets were sold to Tiny Ltd and most of the team joined Temporal.

Sources

Keycard: docs.keycard.ai/reference/security-architecture; fintech.global/2025/10/23/keycard-secures-38m-to-advance-ai-agent-security
Arcade.dev: docs.arcade.dev/en/build/tool-calling/custom-apps/auth-tool-calling; fenwick.com/insights/experience/fenwick-represents-arcade-60m-series-funding
Okta / Auth0: auth0.com/blog/auth0-for-ai-agents-generally-available; developer.okta.com/docs/concepts/xaa
AWS Bedrock AgentCore Policy: docs.aws.amazon.com/bedrock-agentcore/latest/devguide/policy-core-concepts.html; docs.aws.amazon.com/bedrock-agentcore/latest/devguide/policy-temporal.html
NVIDIA OpenShell + Claude Managed Agents: claude.com/blog/giving-companies-more-control-over-their-ai-agents-with-nvidia
Tigera Lynx: tigera.io/blog/nvidia-openshell-secures-the-agent-who-governs-the-fleet; opensourceforu.com/?p=99813
Outerlimit: outerlimit.com/platform/enforce; securitybrief.co.uk/story/outerlimit-raises-usd-16-million-to-secure-ai-agents
Rein Security: pulse2.com/rein-security-raises-25-million-series-a-to-secure-enterprise-ai-agents
Descope: docs.descope.com/agentic-identity-hub/core-components/agents; descope.com/press-release/agentic-identity-hub
Permit.io: docs.permit.io/permit-mcp-gateway/architecture; docs.permit.io/permit-mcp-gateway/human-in-the-loop
Kontext Security: github.com/kontext-security/kontext (commit 95d4da9); securityweek.com/kontext-security-emerges-with-4-million-for-ai-agent-runtime-controls
Authrizd, RuntimeAI: authrizd.ai/technical; runtimeai.io/products
Cerbos, AuthZed, Pomerium: github.com/cerbos/cerbos; github.com/authzed/spicedb; pomerium.com/docs/capabilities/mcp/limit-mcp-tools
Strata Identity: rubrik.com/company/newsroom/press-releases/26/rubrik-advances-identity-resilience-through-strata-acquisition-and-identity-roll-forward-innovation
Oso: osohq.com/post/oso-has-been-acquired-by-tiny
Azure attack (Storm-3168): microsoft.com/en-us/security/blog/2026/09/25/storm-3168-agentic-driven-cloud-attacks-using-compromised-service-principals
Five Eyes guidance on agentic AI (May 2026): crowell.com/en/insights/client-alerts/american-and-allied-cyber-agencies-issue-first-joint-guidance-on-securing-agentic-ai
Cisco on AI agent safety (October 2026); Cisco and Astrix (May 2026): finance.yahoo.com/video/ai-agent-safety-will-be-the-issue-of-our-time-cisco-president-143905775.html; securityweek.com/cisco-moves-to-acquire-astrix-security-to-tackle-non-human-identity-risks
Elliptic-curve research (October 2026): openai.com/index/building-scientific-understanding-with-academic-communities; github.com/bochum-crypto/subexponential-ecdl
Identity-based encryption standards: NCSC MIKEY-SAKKE FAQ (2016); RFC 6508; ETSI TS 133 180; ISO/IEC 18033-5:2015; ETSI TR 103 618 (2019). Also RFC 9396; NIST FIPS 203 and 204
Cyera: en.globes.co.il/en/article-cyera-raising-funds-at-20b-valuation-report-1001561449; cyera.com/press-releases; oasis.security/blog/introducing-oasis-agentic-access-management; thenextweb.com/news/cyera-400m-goldman-sachs-series-g-extension
AxonFlow: github.com/getaxonflow (release v10.2.0; Claude Code plugin); getaxonflow.com/case-studies/bukuwarung

SHOGI SYSTEMS | THE AUTHORITY LAYER

The access layer is crowded. The authority layer is not.

Twenty existing platforms, mapped by what their cryptography proves and where they enforce



Reading the map

- In production
- Designed or claimed, not in production
- SHOGI, in development

Crowded bottom half: identity, policy and gateways prove who is calling, or that a log was not changed.

Top left: three cryptographic designs or claims, none proven in production.

SHOGI is alone where authority is proven and enforced at the point of action.

Source: The Authority Layer, SHOGI Systems, November 2, 2026. Positions reflect public documentation only; SHOGI reflects its design.

SHOGI SYSTEMS | THE AUTHORITY LAYER

One compromised agent, or one person, should never be enough.

The same twenty platforms, mapped by how narrowly an action is authorized and who must agree



Reading the map

- In production
- Designed or claimed, not in production
- SHOGI, in development

Eleven are broad and automated. Four check each call automatically.

Five pause for one person, who becomes a single point of failure.

None requires several independent approvals in production. AxonFlow has started on it.

Only SHOGI requires a one-time permit approved by a quorum for every critical action: no single agent or person can push one through.

Source: The Authority Layer, SHOGI Systems, November 2, 2026. Positions reflect public documentation only; SHOGI reflects its design.